

IoTPentBench: A Network-Scale Benchmark for LLM-Based IoT Penetration Testing

Anonymous Authors

Paper #XXXX, Anonymized for Double-Blind Review

Abstract—[TODO: Abstract — write last. Lead with the network-scale gap, not the harness.]

Index Terms—IoT security, penetration testing, large language models, autonomous agents, benchmark, cyber-physical systems

1. Introduction

[TODO: Write ¶1: IoT attacks are network-scale; defenders pentest with single-host tooling.]

[TODO: Write ¶2: LLM agents are progressing fast, cite 4-5 recent peers.]

[TODO: Write ¶3: formalize the structural gap. Cite 5-6 benchmarks. Keep it sharp.]

Contributions. We bridge this gap with:

- **IoTPentBench** — the first reproducible benchmark for LLM pentest agents at *network scale*: 5 architectural patterns, 7 scenarios, 8 protocols, 87 vulnerabilities deterministically injected via Ansible, with ground truth mapped to OWASP IoT Top 10 and MITRE ATT&CK for ICS.
- **MESHSCOUT** — a network-native harness with *discovery* → *per-device parallel triage* → *per-vulnerability verification* → *network-level report*, designed to keep LLM context tractable at /24 scale.
- **Evidence-level scoring (L1/L2/L3)** — a severity-weighted evaluator that distinguishes detection, exploitation, and data exfiltration, beyond the binary flag-capture of single-host benchmarks.
- **Empirical study** — using a single shared LLM (MiniMax-M2.7) to isolate architectural contributions, we compare MESHSCOUT against three categories of baselines: a non-LLM scanner (Nmap NSE), single-agent LLM (PentestGPT), and multi-agent LLM systems with and without graph guidance (CAI in two scope variants, VulnBot with task-graph). We additionally report two ablations of our pipeline that isolate the individual contributions of the infrastructure-graph phase and the multi-hop intrusion phase.

[TODO: Write ¶5: one teaser sentence per RQ. e.g. "Best LLM reaches XX% recall at /24 scale; scanners reach YY%; but only ZZ% of LLM findings achieve L2+."]

Paper organization. Section 2 surveys LLM pentest agents and benchmarks. Section 3 formalizes the threat model. Section 4 presents IoTPentBench. Section 5 describes the harness architecture. Section 6 and Section 7 present experiments and results. Section 8 and Section 9 discuss limitations and ethics. Artifacts are released at ANONYMIZED_URL.

2. Related Work

This section reviews prior work relevant to our approach. We first survey LLM-driven penetration-testing agents and the benchmarks used to evaluate them, then discuss the classical attack-graph literature that informs MESHSCOUT's topology-aware design, and finally review IoT firmware analysis and the security frameworks underlying our scenario taxonomy.

2.1. LLM Penetration-Testing Agents

Happe and Cito [1] reported the first study of LLM-driven penetration testing, showing that an LLM with shell access could autonomously compromise a small set of vulnerable VMs; the HACKINGBUDDYGPT project [2] extended this line into an open-source agent. PENTESTGPT [3] reused the same single-agent plus human-in-the-loop blueprint and was the first to scale evaluation, using 182 hand-curated HackTheBox and VulnHub sub-tasks; it established the recurring failure mode that motivates all subsequent work, namely state collapse over long command sequences. AUTOATTACKER [4] and NEBULA [5] are contemporaneous single-agent variants oriented towards offensive operations and operator assistance.

Multi-agent decomposition addresses context exhaustion through role specialisation. VULNBOT [6] runs three phase-specific agents (reconnaissance, scanning, exploitation) coordinated by a Penetration Task Graph and reports 30.3% task completion on AUTOPENBENCH, a 21-point gain over single-agent baselines. AUTOPENTESTER [7] adds Retrieval-Augmented Generation and a Pentest Tree across five cooperating roles, reporting a 27-point gain over PENTESTGPT. PENTESTAGENT [8] (ASIA CCS'25) and the production-grade CAI [9], from the PENTESTGPT authors, follow the same pipeline-with-RAG blueprint.

A complementary thread improves the underlying agent rather than its orchestration. PENTEST-R1 [10] applies two-stage reinforcement learning on reasoning traces, and xOF-

FENSE [11] fine-tunes Qwen3-32B on offensive-knowledge corpora, reporting 79% sub-task completion on CTF targets. ENIGMA [12] pairs an LLM with interactive shell tools on CTF challenges, and IRIS [13] couples an LLM with static analysis to surface code-level vulnerabilities. For long-running sessions, CHAP [14] relays distilled context across successive agent instances and AWE [15] adds persistent memory plus browser-backed verification. CHECKMATE [16] couples an LLM with a symbolic planner to make action selection deterministic, improving benchmark success by over 20 points.

Despite the diversity of these designs, every system above is evaluated on a *single host*: a HackTheBox machine, a VulnHub VM, or a CTF container. None reasons over a multi-device IP network with firewall-enforced reachability or IoT-specific application protocols.

2.2. LLM Cybersecurity Benchmarks

The dominant academic corpus is AUTOPEN-BENCH [17], consisting of VulnHub-derived single-VM challenges with scripted flag checkers; it is used by VULNBOT and PENTEST-R1. Recent benchmarks broaden the scope without changing the unit of evaluation: NYU CTF BENCH [18] contributes 200 Jeopardy challenges with reproducible containers, CYBENCH [19] co-evaluates capability and risk on a curated CTF subset, and CYBERSECEVAL [20] measures cybersecurity risk and helpfulness across a large prompt corpus. CAIBENCH [21] consolidates the field into a meta-benchmark spanning Jeopardy CTFs, attack-and-defense ranges, and knowledge tests; its RCTF2 subset adds 27 robotics challenges — the first cyber-physical inclusion in this benchmark family — yet each challenge remains an isolated container.

A separate community ships reproducible multi-host labs. LUDUS [22] and the GOAD [23] Active-Directory scenario deliver Proxmox+Ansible deployments of representative enterprise networks, methodologically the closest precedent to our infrastructure. They ship no ground-truth schema and no scoring procedure, and have not been used to evaluate LLM agents in the published literature. No prior benchmark therefore offers the combination on which IoT-PentBench is built: network-scale evaluation, IoT/OT application protocols (MQTT, Modbus, LoRaWAN, CoAP), per-vulnerability machine-readable ground truth, and evidence-level granularity beyond binary flag capture (Table 1).

2.3. Attack Graphs and Network Reachability

Representing a network as a graph for security analysis predates deep learning by decades. Phillips and Swiler [24] introduced graph-based network-vulnerability analysis, Sheyner *et al.* [25] developed automated attack-graph generation via model checking, and MULVAL [26] cast the problem as Datalog inference, becoming the de facto reference for attack-graph research. These systems compute reachable attacker states from a *static* encoding of host configurations and exposed services. We borrow

their reachability abstraction — nodes are devices, edges are firewall-permitted protocol flows — but MESHSCOUT operates on a *live* network: edges are validated by actual probes, the LLM proposes the next action, and the graph is updated as new hosts are discovered. Classical attack-graph generation is therefore complementary rather than competing.

2.4. IoT Security Context

Scenario design and ground-truth taxonomy follow three references: the OWASP IoT Top 10 [27] for vulnerability categories, MITRE ATT&CK for ICS [28] for OT-tactic annotations, and NIST SP 800-82 Rev. 3 [29] for segmentation conventions, notably the IT/OT split underpinning pattern P3 (§4.2). On the empirical side, the closest IoT security work targets firmware rather than deployed networks: Costin *et al.* [30] performed the first large-scale firmware study, FIRMADYNE [31] introduced dynamic firmware emulation for vulnerability discovery, and FIRMAE [32] substantially improved emulation success rates. OWASP IOTGOAT [33] provides a deliberately insecure firmware image used as a teaching artefact. These projects produce vulnerable images; IoT-PentBench produces a deployed multi-device network with firewall-enforced reachability, the dimension along which we evaluate MESHSCOUT.

3. Threat Model and Problem Definition

Network-scale pentest. We define *network-scale penetration testing* as the task of: (i) discovering all live hosts in a target IP range, (ii) enumerating services and identifying vulnerabilities on each host, (iii) producing evidence of exploitability where feasible, and (iv) aggregating findings at the network level. The *unit of evaluation* is the network, not an individual host.

Attacker model. The attacker holds network-layer access to the target subnet (post-ingress, pre-compromise). They hold no prior credentials, no privileged vantage, and no knowledge of deployed software beyond what is observable through scans. The attacker has access to standard offensive tooling (Nmap, MQTT clients, SSH, ...) and an LLM with tool-calling. Time and budget are bounded.

Goals.

- 1) Enumerate all devices reachable from the entry point.
- 2) For each device, identify vulnerabilities from misconfiguration, weak credentials, information disclosure, insecure protocols, and known CVEs.
- 3) For each finding, attempt to produce *evidence* at one of three levels: L1 detection (port open, banner leak), L2 exploitation (authenticated session, protocol handshake), L3 data exfiltration (configuration files, credentials, sensor data).
- 4) Report findings at the network level (severity-ranked, deduplicated).

TABLE 1. COMPARISON WITH PRIOR LLM PENTEST SYSTEMS AND BENCHMARKS. *Scope* IS THE UNIT OF EVALUATION (*single-host* = ONE VM OR CONTAINER; *network* = MULTI-DEVICE IP RANGE). *Topologies* IS THE DIVERSITY OF TARGET NETWORK ARCHITECTURES.

System	Scope	Protocols	Topologies	IoT-specific	Scoring	Ground truth	Reproducible
Happe & Cito [1]	single-host	Web/OS	n/a	no	flag	hand-curated	partial
PentestGPT [3]	single-host	Web/OS	n/a	no	sub-task	hand-curated	partial
VulnBot [6]	single-host	Web/OS	n/a	no	task-graph	AUTOPENBENCH	yes
AutoPentester [7]	single-host	Web/OS	n/a	no	sub-task	HTB + custom	yes
PentestAgent [8]	single-host	Web/OS	n/a	no	composite	custom	partial
CAI [9]	single-host	Web/CTF	n/a	no	CTF	CAIBench	yes
Pentest-R1 [10]	single-host	Web/OS	n/a	no	sub-task	AUTOPENBENCH	partial
xOffense [11]	single-host	Web/CTF	n/a	no	sub-task	CTF	partial
EnIGMA [12]	single-host	Web/CTF	n/a	no	flag	NYU CTF	yes
CheckMate [16]	single-host	Web/OS	n/a	no	sub-task	custom	partial
AUTOPENBENCH [17]	single-host	Web/OS	n/a	no	flag	yes	yes
NYU CTF Bench [18]	single-host	CTF	n/a	no	flag	yes	yes
Cybench [19]	single-host	CTF	n/a	no	flag	yes	yes
CyberSecEval [20]	prompt-set	mixed	n/a	no	risk score	yes	yes
CAIBench / RCTF2 [21]	single-host	mixed	n/a	27 CPS	CTF	yes	yes
Ludus + GOAD [22], [23]	multi-host	AD / Win	AD only	no	(training)	no	yes
IoTPentBench (ours)	network	MQTT, Modbus, LoRa, ...	5 patterns	yes	L1/L2/L3 + wt.	per-vuln YAML	yes

Scope and non-goals. We evaluate at the *IP/application layer* on deployed IoT networks. We do *not* evaluate: (a) physical/side-channel attacks on hardware, (b) RF-layer attacks requiring SDR captures, (c) firmware extraction from physical devices, (d) post-exploitation persistence or impact, (e) insider threats or supply-chain compromise. Hardware-assisted attacks are left as future work (§10).

4. IoTPentBench: A Network-Scale Benchmark

4.1. Design principles

IoTPentBench is built around four principles derived from the gap analysis in §2:

- Reproducibility.** The entire infrastructure deploys from Ansible playbooks against a Proxmox cluster; every vulnerability is injected idempotently with no manual post-deployment steps.
- Architectural diversity.** We synthesize five network patterns (§1) with distinct attacker reachability graphs enforced by OpenWrt firewall rules, not just by service labels.
- Protocol coverage.** Each scenario mixes multiple IoT and OT protocols (MQTT, Modbus, LoRaWAN gateways, CoAP, SNMP, plus standard SSH/HTTP/FTP/Telnet) to exercise protocol-specific reasoning.
- Ground-truth injectability.** Every injected vulnerability is documented in YAML with device, IP, role, severity, category, OWASP IoT and MITRE ICS mappings, optional CVE, indicators (expected detection strings), and a verification command.

4.2. Five architectural patterns

[TODO: Write §4.2 prose — for each of the 5 patterns: (i) one sentence naming the pattern, (ii) the real-world IoT deployment it models, (iii) the firewall reachability invariant.]

Concretely:

- 1) **Flat.** All devices are directly reachable from the entry point; no pivot is required.
- 2) **Gateway-centric.** A single IoT gateway mediates all access to back-end services; inter-service traffic is blocked without gateway compromise.
- 3) **Segmented IT/OT.** Three zones (admin/IT/OT) with access-control lists; the OT zone (Modbus PLCs, HMI) is reachable only via a jump host.
- 4) **Hub-centric star.** A central home-automation hub is the only device that can reach MQTT broker, DB, and camera; peripherals cannot talk to each other directly.
- 5) **Edge-cloud pivot.** Edge subnet and cloud subnet are separated; only the edge gateway bridges them, making cloud-side services unreachable without edge compromise.

4.3. Deterministic vulnerability injection

Vulnerabilities are injected by a single Ansible role (04_inject_vulns.yml) parameterized by `scenario_id`. Each scenario YAML lists `router_vulns` (OpenWrt-level: Telnet, WAN admin, FTP) and per-service role definitions that trigger role-specific injections: *e.g.*, the `mqtt_broker` role sets `allow_anonymous` true and binds on 0.0.0.0:1883; the `ssh_server` role installs weak credentials (`admin:admin`) and enables `PasswordAuthentication`; the `db_server` role starts MariaDB with a passwordless root. Category distribution is shown in §2.

Fig. 1 placeholder: five topology diagrams side-by-side.
(Flat) — (Gateway-centric) — (Segmented IT/OT) — (Hub-star) — (Edge-cloud)

Figure 1. The five architectural patterns in IoTpentBench. Each pattern is enforced by OpenWrt firewall rules that constrain the attacker’s reachability from the entry point, not by service labels alone.

TABLE 2. VULNERABILITY CATEGORIES AND SEVERITY DISTRIBUTION ACROSS IOTPENTBENCH (87 INJECTED INSTANCES).

Category	Count	CRIT	HIGH	MED/LOW
misconfiguration	17	3	8	6
data_exposure	12	0	7	5
default_credentials	11	6	5	0
info_disclosure	10	0	0	10
no_authentication	9	2	4	3
insecure_update	4	0	0	4
weak_crypto	3	0	0	3
missing_header	3	0	0	3
known_cve	3	0	3	0
code_injection	3	2	1	0
privilege_escalation	2	0	2	0
Total	87	13	30	44

```
- id: V1
  device: s1-mqtt
  ip: "192.168.100.11"
  role: mqtt_broker
  title: "MQTT without authentication"
  severity: high
  category: misconfiguration
  owasp_iot: "I1 / I9"
  mitre_ics: "Initial Access, Collection"
  cve: null
  indicators:
    - "allow_anonymous true"
    - "Anonymous clients connected"
  verification: "mosquitto_sub -h 192.168.100.11 -t '#' -v"
```

Figure 2. Ground-truth entry example.

4.4. Ground-truth schema

Each scenario ships a YAML file describing every injected vulnerability. An entry example (from `scenario_1.yaml`):

4.5. Evaluator

The evaluator (§ 5.6 integrates results into Phase 5) matches each LLM finding against ground truth in four passes, in order: (1) exact CVE match; (2) (ip, vuln_type); (3) (ip, category); (4) (ip, severity). Severity weights are $w = \{\text{critical}=4, \text{high}=3, \text{medium}=2, \text{low}=1\}$; loose category matches apply $0.5\times$, severity mismatches $0.75\times$. The composite score is the weighted True-Positive sum normalized by $\sum_{g \in \text{GT}} w(g)$.

Evidence levels. Beyond match accuracy, the evaluator tracks, per True Positive, the deepest *evidence level* achieved by the harness: L1 (detection — banner or port), L2 (exploitation — authenticated action),

L3 (data exfiltration — sensitive content retrieved). We report $Lk\text{-coverage} = |\{t \in \text{TP} : \text{level}(t) \geq k\}|/|\text{TP}|$. This metric cannot be computed on single-host binary-flag benchmarks and captures a dimension of pentest quality that existing LLM-pentest evaluations conflate.

5. MESHSCOUT: A Network-Native LLM Pentest Harness

5.1. Overview

The harness decomposes network pentest into five phases with typed deliverables. The key design decision is *parallelization at the device and vulnerability granularity*: rather than a single monolithic agent that must hold the full network state, we spawn dedicated micro-agents per target with constrained tool sets, then aggregate deterministically. This keeps LLM context bounded at $O(1)$ per device regardless of network size.

5.2. Phase 1: Topology prior

Phase 1 loads a graph model of the target network as a directed graph $G = (V, E)$ where V are devices and E are links labeled by protocol. Two modes are supported: (i) *scenario mode* loads the predefined topology from the benchmark YAML; (ii) *discovery mode* starts from an empty graph and instructs Phase 2 to populate V via network scans. The output is a markdown deliverable enumerating attack surface and preliminary pivot candidates (betweenness-centrality ranked).

5.3. Phase 2: Full-network discovery

Phase 2 executes active reconnaissance across the target /24: `nmap_discovery` (host discovery), `nmap` (service/version), `arp_scan` (L2 + vendor), and protocol-specific probes (`mqtt_listen` subscribing to # on discovered brokers, `ssh_audit` against SSH servers, `curl_headers` on HTTP). Unlike single-host harnesses, the agent is not given a target IP — it must derive one from the CIDR scope.

5.4. Phase 3: Per-device parallel triage

Phase 3 runs in three sub-phases:

- 1) **Deterministic scanner pass.** Ansible-driven subprocess wrappers (`ssh_audit`, `curl_headers`, `mqtt_listen`, `modbus_scan`, ...) run in parallel

Fig. 2 placeholder: 5-phase pipeline diagram.
Phase 1 (topology prior) → Phase 2 (discovery) → Phase 3 (per-device triage) → Phase 4 (per-vuln verification) → Phase 5 (network report).

Figure 3. MESHSOUT pipeline. Solid arrows are deliverables; dashed boxes indicate LLM agent calls; Phase 3 and Phase 4 spawn micro-agents in parallel (§5.4–§5.5).

Fig. 3 placeholder: Phase 3 parallelization — n devices → n LLM agents in parallel → deterministic merge.

Figure 4. Per-device parallelization. Each micro-agent receives only its device’s scan output, avoiding the context explosion of a monolithic network-level agent.

per discovered device and dump structured scan results to disk.

- 2) **Per-device LLM triage.** For each device, a dedicated LLM agent receives only that device’s Phase 2+3a outputs and has access to a minimal tool set: `http_get` (follow-up file retrieval) and `cve_search` (CPE-to-CVE lookup via NVD). The agent produces a JSON vuln list scoped to that device.
- 3) **Deterministic merge.** Per-device outputs are concatenated, canonicalized through the taxonomy (§5.7), and deduplicated via severity-aware rules.

This decomposition is essential for network scale: a monolithic agent on a 15-device scenario would need to fit ~30–80k tokens of scan output into a single context and reason about cross-device relationships simultaneously. Per-device micro-agents flatten this to a constant size, at the cost of losing some cross-device inference capability, which is recovered in Phase 5.

5.5. Phase 4: Per-vulnerability verification with evidence levels

Phase 4 spawns one micro-agent per Phase 3 finding (in parallel, with a concurrency cap). Each micro-agent receives the finding (IP, type, port, details) and full operational tool access: `ssh_login`, `mysql_query`, `mqtt_listen`, `http_get`, `ftp_list`, `modbus_scan`, `telnet_connect`. The agent is prompted to produce the deepest evidence level it can (L1–L3, §4.5) within a bounded turn budget. Outputs include: `status` (CONFIRMED, NOT_EXPLOITABLE, ERROR), `evidence_level`, `command executed`, `stdout excerpts`, and `data_extracted` (when L3).

Findings in CONFIRMED state from Phase 3 override Phase 4 ERROR/FAILED statuses (rationale: Phase 3 is hypothesis generation with directly-observed indicators; Phase 4 crashes should not erase directly-observed findings).

5.6. Phase 5: Network-level report

Phase 5 consumes all prior deliverables and produces a network-scoped report with: device inventory, per-device

vulnerabilities with severity, attack surface summary, and actionable recommendations ordered by severity × exposure. This is the only phase with cross-device reasoning, deliberately placed at the end to avoid bloating earlier phases.

5.7. Taxonomy and tool abstraction

A single module resolves LLM-produced vulnerability type strings into a canonical taxonomy of 11 categories (§2), filtering noise (e.g., meta-observations) and CONFIG-only findings. Tools are defined declaratively as YAML (schema: name, description, JSON-schema parameters, subprocess command template), so adding a new protocol probe requires no Python code.

5.8. Complexity and cost envelope

For a network with n devices and an average of v findings per device, the harness makes $O(1 + n + nv)$ LLM calls (Phases 1, 3b, 4 respectively), bounded per call by fixed turn budgets (50/50/10/10/20 for phases 1–5). Total cost scales linearly with network size, not combinatorially.

6. Experimental Setup

Infrastructure. All scenarios run on a single Proxmox VE host. Networks are Linux bridges (`vmbr1`, 192.168.100.0/24) with an OpenWrt router as gateway. Per-scenario deployment and teardown are automated via Ansible and complete in under 10 minutes per scenario.

LLM model — fairness lock. To isolate *architectural* contributions from model-specific advantages, all LLM-driven systems compared in this work — our pipeline (and its ablations) and all LLM baselines — run on the *same* general-purpose model: **MiniMax-M2.7** [34]. We deliberately avoid a multi-model matrix: a fair comparison of architectures requires the model variable to be held constant. CAI’s specialised `alias1` model (security-fine-tuned, pay-walled) is intentionally not used — this study addresses the architecture axis, not the model fine-tuning axis. The choice of MiniMax-M2.7 also keeps the combined API cost at zero under our research subscription, enabling the full $k=3$ -run variance estimation across all systems and scenarios without budget trade-offs.

Baselines. We compare our pipeline against three categories of systems, each occupying a distinct architectural slot:

- **Single-agent LLM** — PENTESTGPT [3] in autonomous mode, the canonical USENIX’24 reference.

- **Multi-agent LLM (no infrastructure graph)** — CAI [9]. We test two variants: *A* per-IP (one independent CAI session per ground-truth IP, no shared state) and *B* CIDR-scope (a single CAI session given the full /24, allowed to discover hosts and pivot via LinuxCmd primitives). Both variants receive the same total turn budget as our pipeline.
- **Multi-agent LLM with task-graph** — VULNBOT [6], the most structurally similar competitor. VulnBot’s Penetration Task Graph (PTG) is a directed acyclic graph of *tasks* (action dependencies), to be contrasted with our infrastructure graph of *devices* (firewall-enforced reachability). Both are “graph-based”, but encode orthogonal abstractions.
- **Non-LLM scanner** — Nmap with all NSE scripts (`--script=default,vuln,auth`) against the same target CIDR. Scanner outputs are converted to the evaluator’s JSON schema.

Internal ablations. Two ablations isolate the contribution of our two key design choices: (i) MESHSCOUT-w/o-Phase 1 disables graph analysis (`--phases 2 3 4 5 6`), testing the value of infrastructure-graph guidance; (ii) MESHSCOUT-w/o-Phase 5 disables the intrusion phase (`--phases 1 2 3 4 6`), testing the value of structured multi-hop lateral movement.

Variance. Each (system, scenario) cell is run $k=3$ times. We report mean, standard deviation, and 95% bootstrap confidence intervals. Pairwise differences (our pipeline vs each baseline) are tested with Mann–Whitney U on F1 and weighted Score per scenario. This protocol exceeds the single-run reporting of PentestGPT [3], AutoPentester [7], AUTOPENBENCH [17], and CAI [9], and matches the multi-run averaging of VulnBot [6] (5-experiment) and Pentest-R1 [10] (pass@3).

Metrics. We report five families of metrics, chosen to (i) align with prior work’s expected vocabulary and (ii) surface architectural differences invisible to host-level benchmarks:

- *Task completion* — Recall against the ground-truth vulnerability set, the network-scale analogue of the sub-task / task completion rate reported by PentestGPT, VulnBot, AutoPentester, Pentest-R1, and xOffense.
- *Quality* — Precision, F1, and weighted Score (CVSS-severity-weighted Recall, §4.5). Score plays the role of AutoPentester’s “vulnerability coverage” but with severity weights.
- *Multi-Hop Reach* — MHR_k for $k \in \{1, 2, 3\}$, the fraction of ground-truth vulnerabilities at $\text{hop_depth} \geq k$ that are reached. No prior LLM-pentest benchmark measures network depth explicitly.
- *Evidence depth* — L_1 (detection), L_2 (exploitation), L_3 (exfiltration) coverage; absent from binary flag-capture benchmarks.

Fig. 4 placeholder: bar chart, F1 and MHR_2 for MESHSCOUT vs CAI (A/B), VulnBot, PentestGPT, Nmap NSE.

Figure 5. Baseline comparison. Scanners capture known CVEs but miss configuration vulnerabilities. Single-host LLM agents (PentestGPT, CAI A1) cannot pivot across firewall boundaries. CAI B (given the full CIDR) and VulnBot (multi-agent + task-graph) approach our pipeline on direct findings but collapse on multi-hop reach (MHR_2 , MHR_3).

Fig. 5 placeholder: heatmap scenario $\times$ model, cell = weighted Score.

Figure 6. Per-scenario Score heatmap. [TODO: Note which pattern consistently breaks LLMs.]

- *Cost envelope* — prompt+completion tokens and monetary cost (USD) from provider billing APIs, plus wall-clock time per scenario, matching CAI [9], AutoPentester, and Pentest-R1.

Compute and budget. All runs consume a combined API budget under \$NNN, bounded by per-phase turn caps and the concurrency limit in Phase 4. [TODO: Insert exact budget after all runs complete (freeze 2 May).]

7. Evaluation

7.1. Overall performance (RQ1)

[TODO: Write §7.1 prose — headline finding: MESHSCOUT achieves $MHR_2 \approx XX\%$, while all baselines collapse to near-zero on multi-hop vulnerabilities. F1 gap to nearest baseline (VulnBot) is ZZ percentage points. Variance across 3 runs stays under YY%.]

7.2. Network-scale vs baselines (RQ4)

[TODO: Write §7.2 — explicit narrative: (i) scanners win on CVE category, (ii) LLMs win on misconfig/default-creds, (iii) single-host LLMs adapted to network fall behind native harness on F1 by ZZ%.]

7.3. Per-architecture analysis (RQ3)

[TODO: Write §7.3 — pattern-by-pattern commentary. Expected finding: gateway-centric and hub-star degrade most, because they require pivot reasoning. Edge-cloud may also degrade on cloud-side service discovery.]

7.4. Evidence-level depth (RQ2)

[TODO: Write §7.4 — expected: detection (L_1) near-parity with scanners, but L_2/L_3 is LLM-only. This is the strongest argument for LLM agents in this context.]

TABLE 3. MAIN RESULTS — TASK COMPLETION (RECALL), PRECISION, F1, WEIGHTED SCORE, MULTI-HOP REACH (MHR_k), TOKENS, AND USD COST; AVERAGED OVER 7 SCENARIOS $\times$ 3 RUNS. ALL LLM SYSTEMS USE MINIMAX-M2.7 (CF. §6); NMAP NSE IS THE NON-LLM LOWER BOUND. *Task Completion* IS THE NETWORK-SCALE ANALOGUE OF THE SUB-TASK COMPLETION RATE REPORTED BY PENTESTGPT, VULNBOT, AUTOPENTESTER, AND XOFFENSE (§7.1). MHR_k MEASURES THE FRACTION OF GROUND-TRUTH VULNERABILITIES AT $HOP_DEPTH \geq k$ THAT WERE REACHED; “–” INDICATES NO GT ENTRIES AT THAT DEPTH. CELLS PRE-FILLED WITH 0.00 ARE ARCHITECTURALLY IMPOSSIBLE: SINGLE-HOST LLMs AND SCANNERS CANNOT PIVOT ACROSS FIREWALL BOUNDARIES BY CONSTRUCTION.

System	Quality				Multi-Hop Reach			Cost	
	TC	Prec.	F1	Score	MHR_1	MHR_2	MHR_3	Tokens (M)	USD
MESHSCOUT (full)	[TODO:]	[TODO:]	[TODO:]	[TODO:]	[TODO:]	[TODO:]	[TODO:]	[TODO:]	[TODO:]
MESHSCOUT w/o Phase 1 (no infra graph)	[TODO:]	[TODO:]	[TODO:]	[TODO:]	[TODO:]	[TODO:]	[TODO:]	[TODO:]	[TODO:]
MESHSCOUT w/o Phase 5 (no intrusion)	[TODO:]	[TODO:]	[TODO:]	[TODO:]	[TODO:]	0.00	0.00	[TODO:]	[TODO:]
CAI [9] (Variant A1, per-IP)	[TODO:]	[TODO:]	[TODO:]	[TODO:]	[TODO:]	0.00	0.00	[TODO:]	[TODO:]
CAI [9] (Variant B, CIDR scope)	[TODO:]	[TODO:]	[TODO:]	[TODO:]	[TODO:]	[TODO:]	[TODO:]	[TODO:]	[TODO:]
VulnBot [6] (multi-agent + PTG)	[TODO:]	[TODO:]	[TODO:]	[TODO:]	[TODO:]	[TODO:]	[TODO:]	[TODO:]	[TODO:]
PentestGPT [3] (single-agent)	[TODO:]	[TODO:]	[TODO:]	[TODO:]	[TODO:]	0.00	0.00	[TODO:]	[TODO:]
Nmap NSE (non-LLM scanner)	[TODO:]	[TODO:]	[TODO:]	[TODO:]	[TODO:]	0.00	0.00	–	–

TC = Task Completion (Recall against GT vuln set). Reported across $k=3$ runs with 95% bootstrap CI and pairwise Mann–Whitney U vs each baseline; this protocol exceeds the single-run reporting of PentestGPT, AutoPentester, AUTOPENBENCH, and CAI, and matches the multi-run averaging of VulnBot and Pentest-R1.

TABLE 4. EVIDENCE-LEVEL COVERAGE. L_k IS THE FRACTION OF TRUE POSITIVES REACHING LEVEL k OR DEEPER. ABSENT FROM BINARY-FLAG BENCHMARKS.

System	L_1 (detect)	L_2 (exploit)	L_3 (exfiltrate)
MESHSCOUT (full)	[TODO:]	[TODO:]	[TODO:]
MESHSCOUT w/o Phase 5	[TODO:]	[TODO:]	0.00
CAI (Variant A1) [9]	[TODO:]	[TODO:]	0.00
CAI (Variant B) [9]	[TODO:]	[TODO:]	[TODO:]
VulnBot [6]	[TODO:]	[TODO:]	[TODO:]
PentestGPT [3]	[TODO:]	[TODO:]	0.00
Nmap NSE	–	–	–

Fig. 6 placeholder: Pareto plot, USD cost vs weighted Score, one point per model.

Figure 7. Cost–performance Pareto. [TODO: Annotate which model dominates on each axis and the efficient frontier.]

7.5. Per-protocol analysis

[TODO: Write §7.5 — 2 paragraphs. Expected: LLMs excel on MQTT/SSH/HTTP (text-rich, well-represented in training), fall short on Modbus/LoRaWAN (binary, sparse training data). Report numbers per protocol.]

7.6. Failure modes and cost/time envelope

[TODO: Write §7.6 — categorize failure modes: (a) hallucinated findings on devices with no services, (b) missed vulnerabilities on binary protocols, (c) over-reporting info-disclosure without severity calibration. Brief quantitative table.]

8. Discussion and Limitations

Where LLMs help, where they don’t. IoTPentBench surfaces an empirical pattern: LLM agents reach high recall on misconfigurations, default credentials, and information disclosure — findings that are textually expressible and well represented in their training corpora. They consistently underperform on binary industrial protocols (Modbus, CoAP, LoRaWAN) and on cross-device inferences that require holding an aggregated network view. This suggests a hybrid deployment model: LLM agents for configuration auditing, traditional scanners for CVE coverage, and human operators for the OT/ICS layer.

Implications for agentic systems. The harness’s per-device and per-vulnerability micro-agent structure is one answer to context scaling, but it cedes cross-device inference. Designs that keep a global aggregator agent in the loop without paying the full network-context cost are an open research direction.

Limitations. (i) *Simulated infrastructure*. IoTPentBench scenarios emulate IoT devices on Linux VMs; real constrained hardware introduces embedded-specific attack surfaces (bootloaders, RTOS, firmware) we do not cover. (ii) *Evaluator leniency*. Bonus categories are accepted as non-FP even without ground-truth matches, which inflates precision. (iii) *Model variance*. With $k=3$ runs per configuration, confidence intervals on F1 remain wide for some scenarios. (iv) *Prompt sensitivity*. We fix a single prompt family across models; per-model prompt tuning would likely raise scores. (v) *Static architectures*. The five patterns are fixed; real deployments evolve over time and include devices not represented here (RFID badge readers, industrial PLCs, medical devices).

9. Ethical Considerations

Isolation. All experiments run on an isolated Proxmox cluster with a dedicated bridge (vmbri1) and no route

to external networks; the OpenWrt gateway’s WAN interface is firewall-isolated. No live hosts, production services, or third-party infrastructure are probed.

Responsible disclosure. IoTPentBench vulnerabilities are deterministically injected into our own lab VMs via Ansible. No 0-days are introduced; all categories (misconfig, weak credentials, Terrapin CVE-2023-48795, outdated Dropbear) correspond to patterns already publicly documented. No responsible-disclosure process is triggered.

Misuse risk of released artifact. The benchmark artifact (Ansible playbooks, ground truth, evaluator) does not embed exploitation payloads for unknown vulnerabilities, and the injected weaknesses are common misconfigurations documented in OWASP IoT Top 10. The harness uses only standard pentest tools already in widespread use. We judge the net societal benefit — reproducible evaluation of LLM pentest agents, and a concrete evaluation surface for future defenders — to exceed the incremental misuse risk.

IRB / human subjects. No human subjects are involved. No personal data is collected or processed.

LLM costs and environmental footprint. All reported experiments together consume an estimated **[TODO: X]** kWh and cost under **[TODO: USD Y]** in API billing.

10. Conclusion

We presented IoTPentBench, the first reproducible benchmark for evaluating LLM penetration testing agents at *network scale* on IoT infrastructure, and MESHSCOUT, a network-native multi-agent harness whose infrastructure-graph guidance and dedicated lateral movement phase keep LLM context tractable while explicitly modelling firewall-enforced reachability. Holding the LLM constant (MiniMax-M2.7) across our pipeline, two ablations, and four external baselines (CAI in two scope variants, VulnBot, PentestGPT, Nmap NSE), we isolate the *architectural* contribution: MESHSCOUT outperforms single-host LLM agents on multi-hop reach (MHR₂, MHR₃) by a margin that collapses near zero for every baseline by construction, and demonstrates that the L1–L3 evidence spectrum exposes architectural differences that binary flag-capture benchmarks conflate.

Future work. (i) *Multi-model robustness* — extending the architectural comparison to recent frontier models (Claude Sonnet 4.6, GPT-5.5, Gemini 3.1 Pro, DeepSeek V4, Qwen 3.6) to confirm model-independence of the architectural advantage; (ii) *Real hardware* — extending to embedded Zigbee, LoRaWAN, and sub-GHz attack surfaces via SDR tooling; (iii) *Defender’s side* — using IoTPentBench as a benchmark for LLM-assisted IoT defense (detection, patching recommendation).

All code, scenarios, ground truth, and run logs are released at ANONYMIZED_URL under an open license.

References

[1] A. Happe and J. Cito, “Getting pwn’d by AI: Penetration testing with large language models,” in *Proc. ACM ESEC/FSE*, 2023.

[2] A. Happe *et al.*, “HackingBuddyGPT: An LLM-driven pentest agent,” 2024. [Online]. Available: <https://github.com/ipa-lab/hackingBuddyGPT>

[3] G. Deng *et al.*, “PentestGPT: Evaluating and harnessing large language models for automated penetration testing,” in *Proc. 33rd USENIX Security Symp.*, 2024.

[4] J. Xu *et al.*, “AutoAttacker: A large language model guided system to implement automatic cyber-attacks,” *arXiv:2403.01038*, 2024.

[5] O. V. Aguinaga, “Nebula: An LLM-powered penetration testing assistant,” *arXiv:2412.00006*, 2024.

[6] H. He *et al.*, “VulnBot: Autonomous penetration testing for a multi-agent collaborative framework,” *arXiv:2501.13411*, 2025.

[7] Y. Ginige *et al.*, “AutoPentester: An LLM agent-based framework for automated pentesting,” *arXiv:2510.05605*, 2025.

[8] X. Shen *et al.*, “PentestAgent: Incorporating LLM agents to automated penetration testing,” in *Proc. 20th ACM Asia Conf. Computer and Communications Security (ASIA CCS)*, 2025. doi: 10.1145/3708821.3733882.

[9] V. Mayoral-Vilches *et al.*, “CAI: An open, bug bounty-ready cybersecurity AI,” *arXiv:2504.06017*, 2025.

[10] Anonymous, “Pentest-R1: Towards autonomous penetration testing reasoning optimized via two-stage reinforcement learning,” *arXiv:2508.07382*, 2025.

[11] Anonymous, “xOffense: An AI-driven autonomous penetration testing framework with offensive knowledge-enhanced LLMs and multi-agent systems,” *arXiv:2509.13021*, 2025.

[12] T. Abramovich *et al.*, “EnIGMA: Interactive tools substantially assist LM agents in solving CTF challenges,” in *Proc. ICML*, 2025.

[13] Z. Li, S. Dutta, and M. Naik, “IRIS: LLM-assisted static analysis for detecting security vulnerabilities,” in *Proc. NDSS*, 2025.

[14] Anonymous, “Context relay for long-running penetration-testing agents,” in *NDSS Workshop on LLM-Assisted Security and Trust Exploration (LAST-X)*, 2026.

[15] Anonymous, “AWE: A memory-augmented multi-agent framework for autonomous web penetration testing,” in *NDSS Workshop on LLM-Assisted Security and Trust Exploration (LAST-X)*, 2026.

[16] Anonymous, “Automated penetration testing with LLM agents and classical planning,” *arXiv:2512.11143*, 2025.

[17] Anonymous, “Towards automated penetration testing: Introducing LLM benchmark, analysis, and improvements,” *arXiv:2410.17141*, 2024.

[18] M. Shao *et al.*, “NYU CTF Bench: A scalable open-source benchmark dataset for evaluating LLMs in offensive security,” in *Proc. NeurIPS Datasets and Benchmarks Track*, 2024.

[19] A. K. Zhang *et al.*, “Cybench: A framework for evaluating cybersecurity capabilities and risk of language models,” *arXiv:2408.08926*, 2024.

[20] M. Bhatt *et al.*, “CyberSecEval 2: A wide-ranging cybersecurity evaluation suite for large language models,” *arXiv:2404.13161*, 2024.

[21] V. Mayoral-Vilches *et al.*, “CAIBench: A meta-benchmark for evaluating cybersecurity AI agents,” *arXiv:2510.24317*, 2025.

[22] Ludus Cloud, “Ludus: Cyber range automation on Proxmox.” [Online]. Available: <https://ludus.cloud>

[23] M3, “GOAD: Game of Active Directory.” [Online]. Available: <https://github.com/Orange-Cyberdefense/GOAD>

[24] C. Phillips and L. P. Swiler, “A graph-based system for network-vulnerability analysis,” in *Proc. New Security Paradigms Workshop*, 1998.

[25] O. Sheyner, J. Haines, S. Jha, R. Lippmann, and J. M. Wing, “Automated generation and analysis of attack graphs,” in *Proc. IEEE Symp. Security and Privacy*, 2002.

- [26] X. Ou, S. Govindavajhala, and A. W. Appel, "MulVAL: A logic-based network security analyzer," in *Proc. USENIX Security Symp.*, 2005.
- [27] OWASP Foundation, "OWASP Internet of Things Top 10," 2024. [Online]. Available: <https://owasp.org/www-project-internet-of-things-top-10/>
- [28] MITRE Corporation, "MITRE ATT&CK for ICS," 2024. [Online]. Available: <https://attack.mitre.org/matrices/ics/>
- [29] K. Stouffer *et al.*, "Guide to operational technology security," National Institute of Standards and Technology, NIST SP 800-82 Rev. 3, 2023.
- [30] A. Costin, J. Zaddach, A. Francillon, and D. Balzarotti, "A large-scale analysis of the security of embedded firmwares," in *Proc. USENIX Security Symp.*, 2014.
- [31] D. D. Chen, M. Egele, M. Woo, and D. Brumley, "Towards automated dynamic analysis for Linux-based embedded firmware," in *Proc. NDSS*, 2016.
- [32] M. Kim, D. Kim, E. Kim, S. Kim, Y. Jang, and Y. Kim, "FirmAE: Towards large-scale emulation of IoT firmware for dynamic analysis," in *Proc. Annual Computer Security Applications Conf. (ACSAC)*, 2020.
- [33] OWASP, "IoTGoat: A deliberately insecure IoT firmware," 2024. [Online]. Available: <https://github.com/OWASP/IoTGoat>
- [34] MiniMax, "MiniMax-M2.7: A general-purpose large language model," 2026. [Online]. Available: <https://www.minimax.io/models/text/m27>